QUIC Logging, debugging and tooling

Robin Marx

UHASSELT EDM

https://quic.edm.uhasselt.be
IETF interim Tokyo – January 2019

QUIC timeline : split per-stream

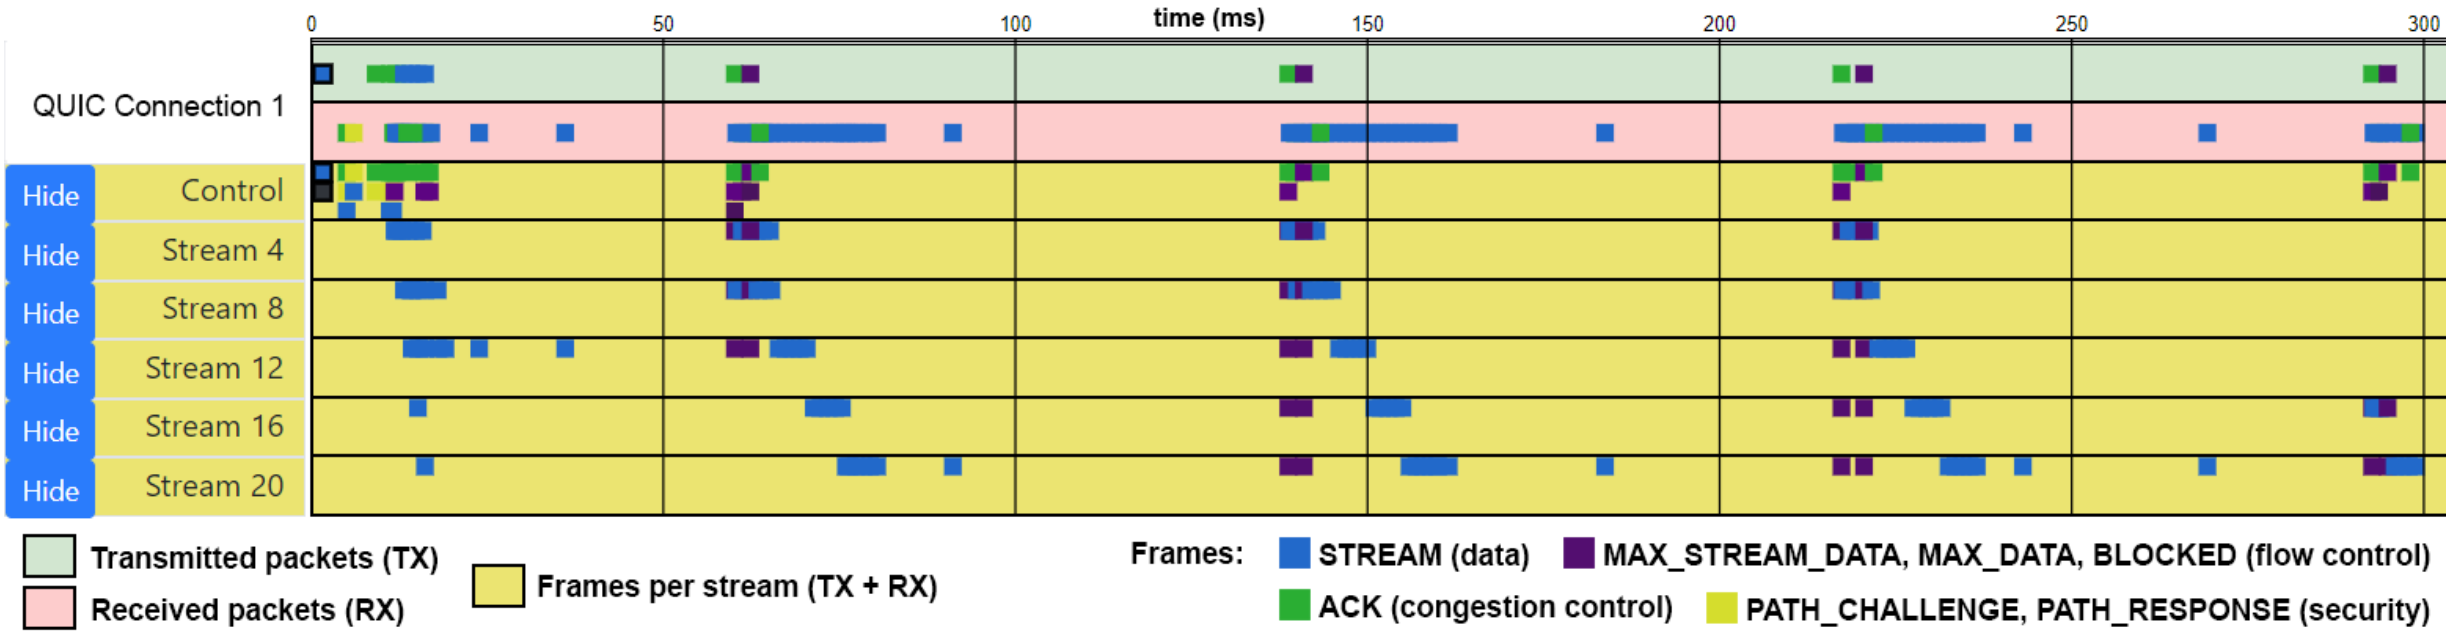

QUIC timeline : compare connections : network drop

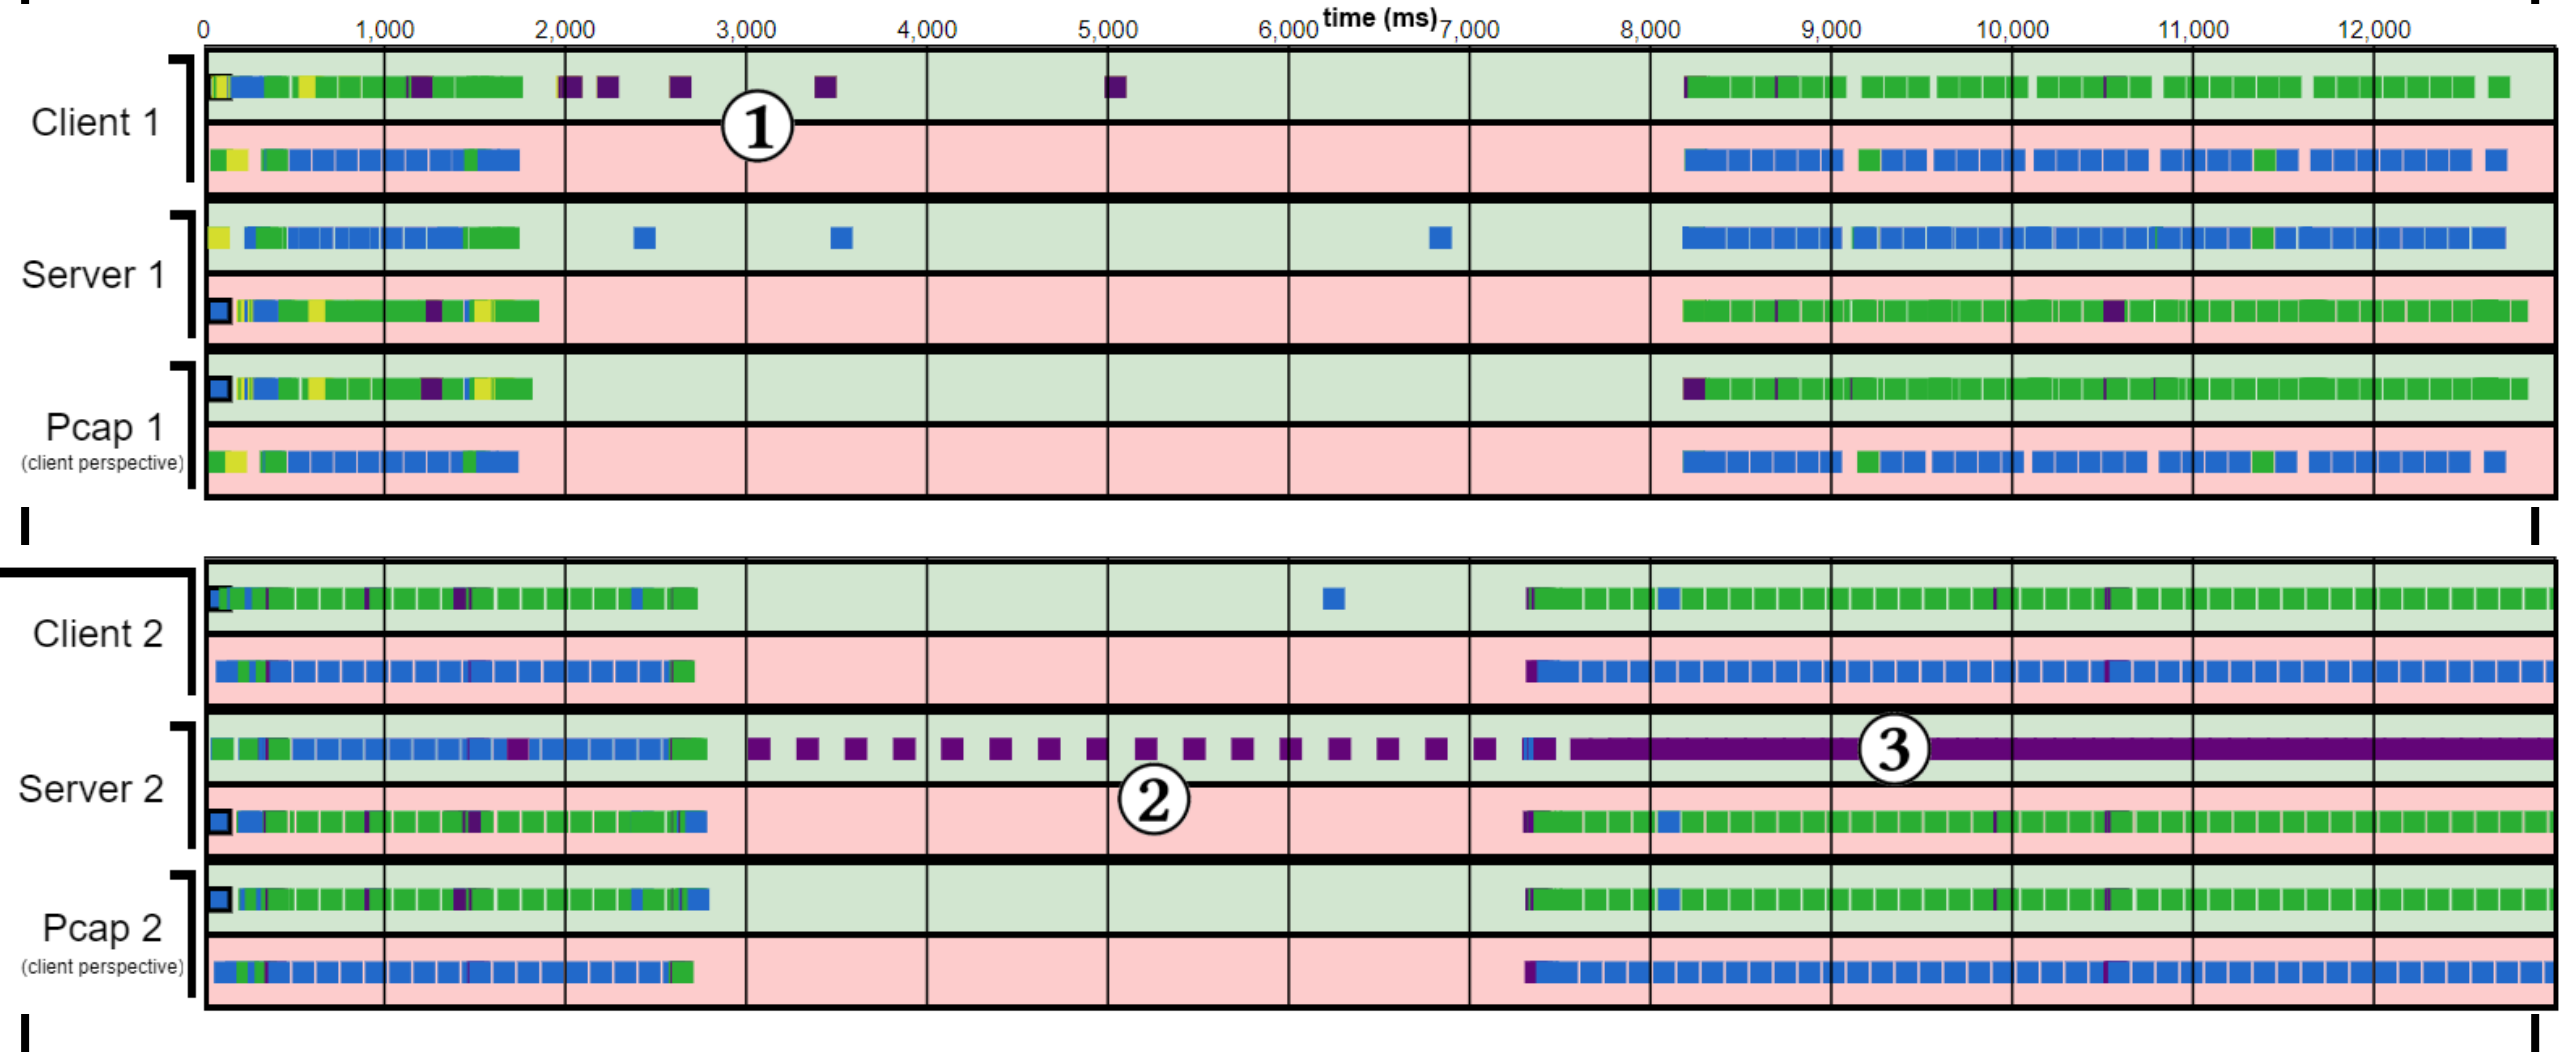

QUIC sequence diagram

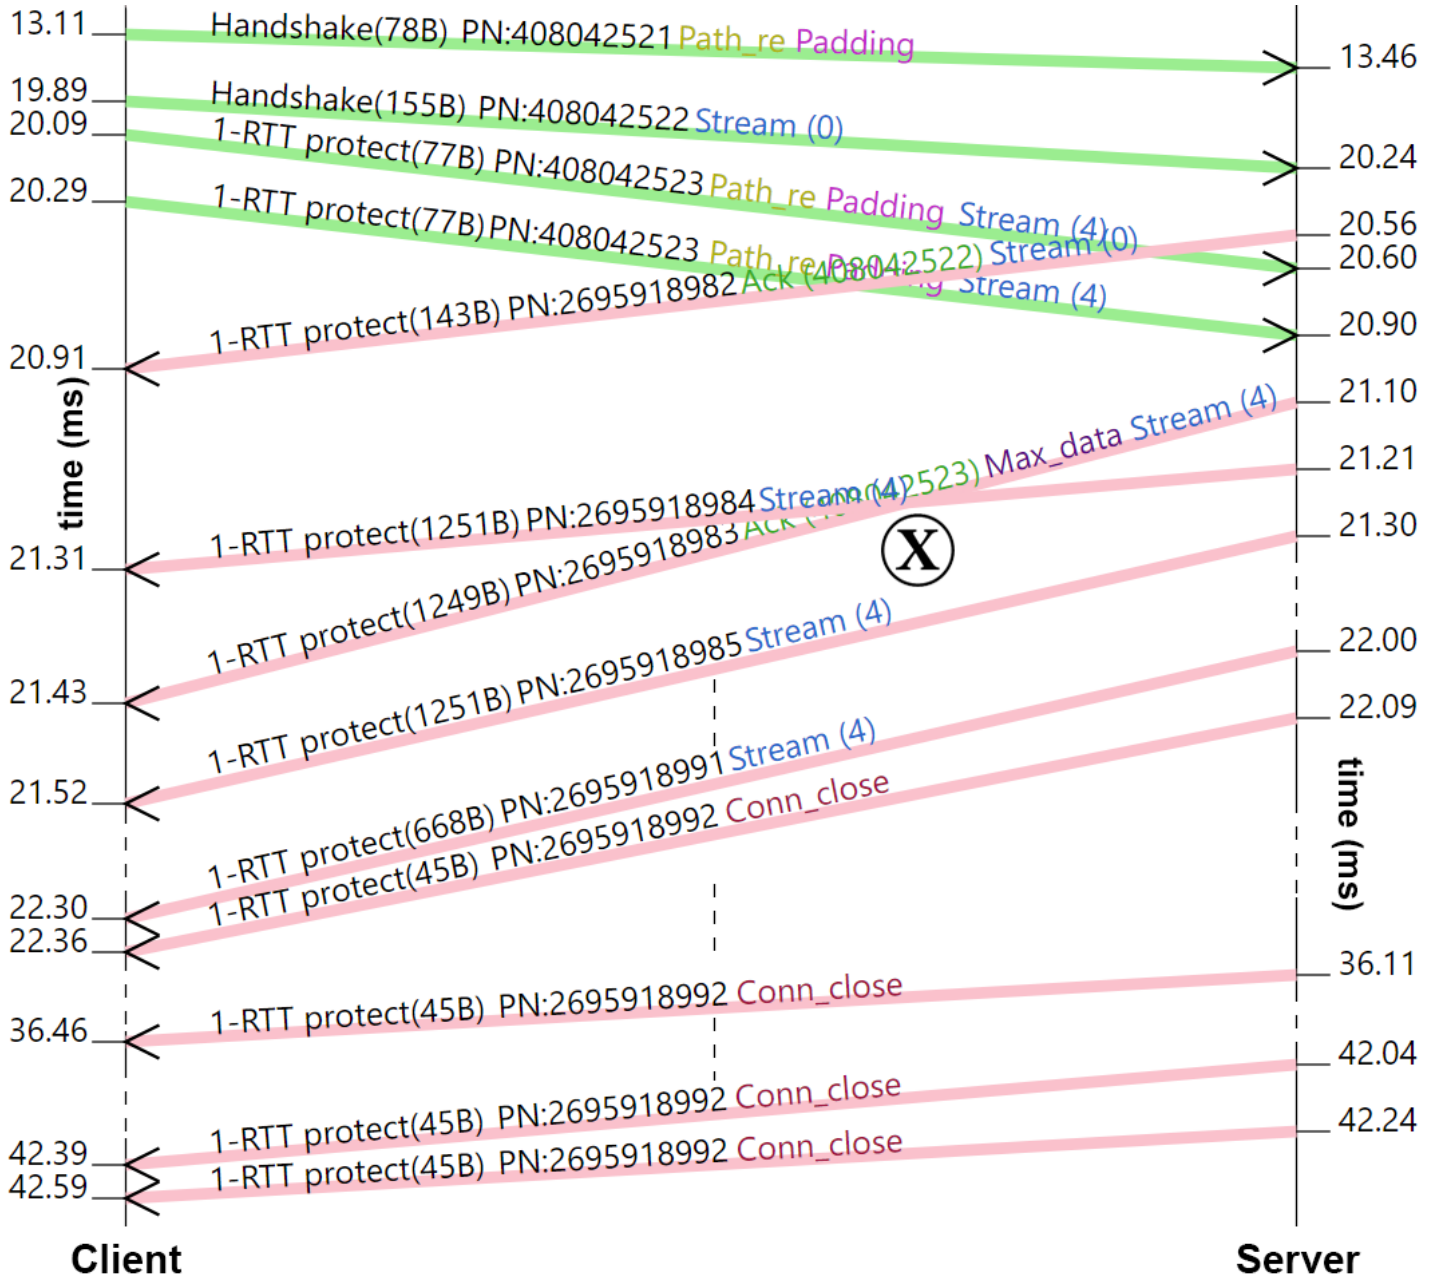

QUIC sequence diagram

Re-ordering

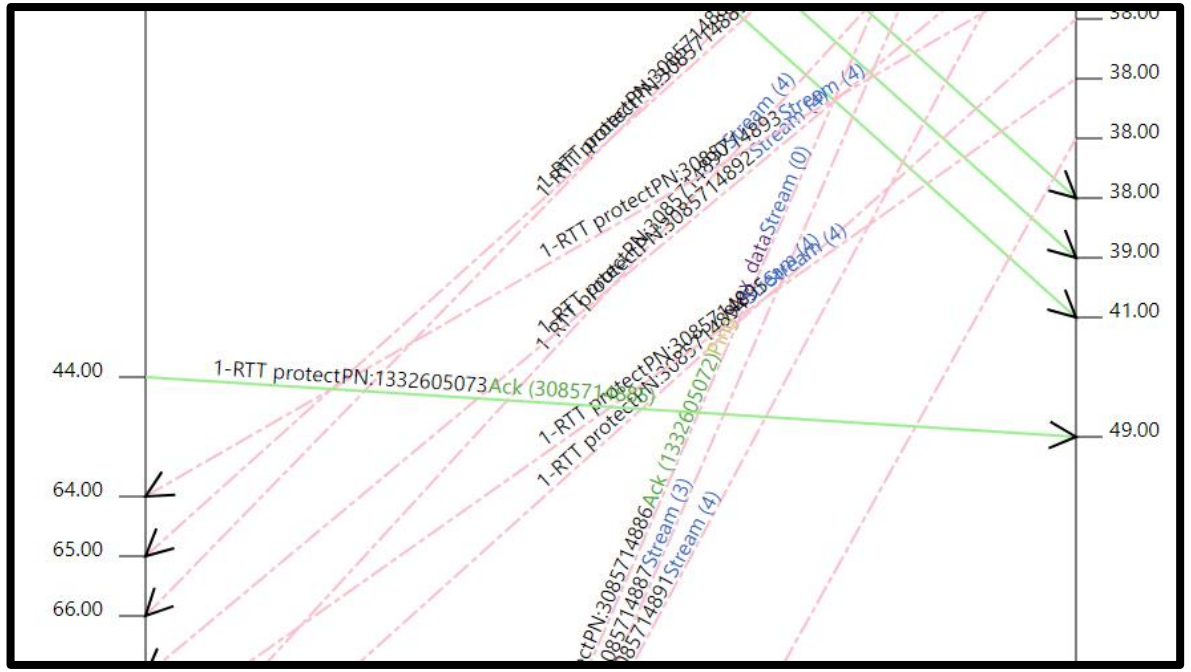

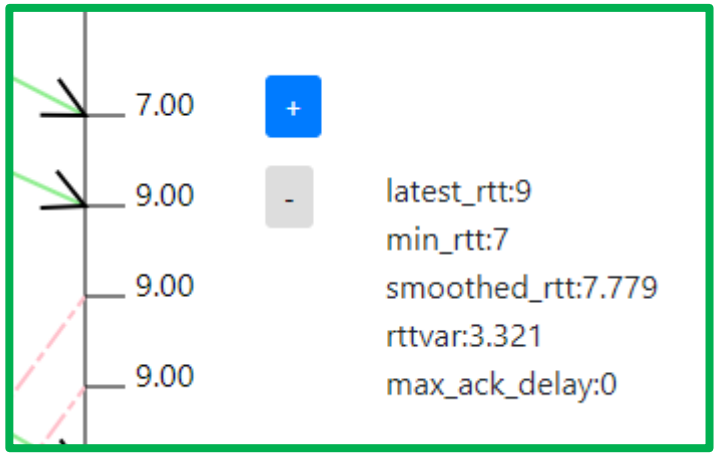

RTT estimates

Loss

Retransmits

QUIC flow and congestion control

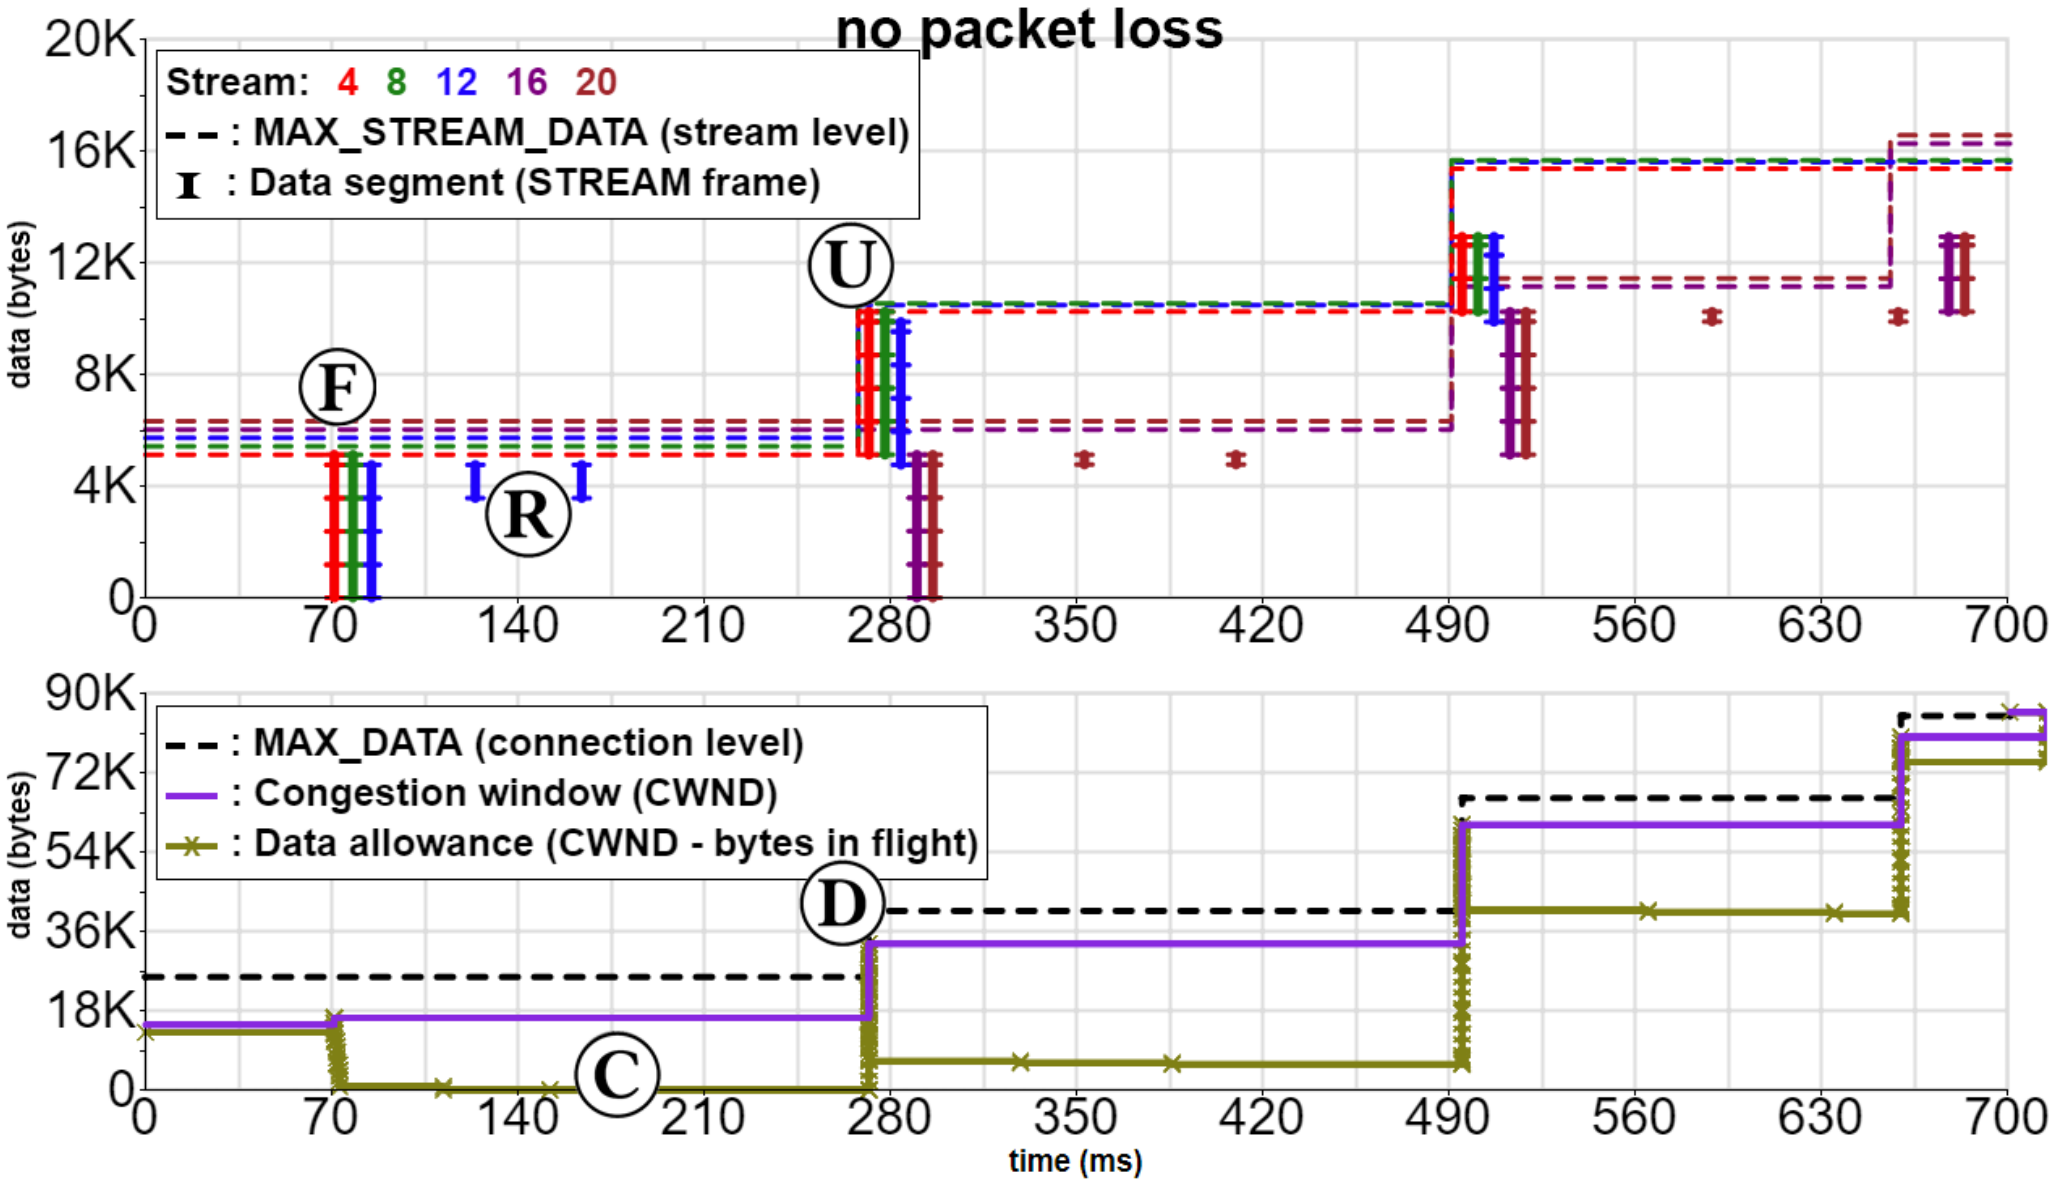

QUIC flow and congestion control

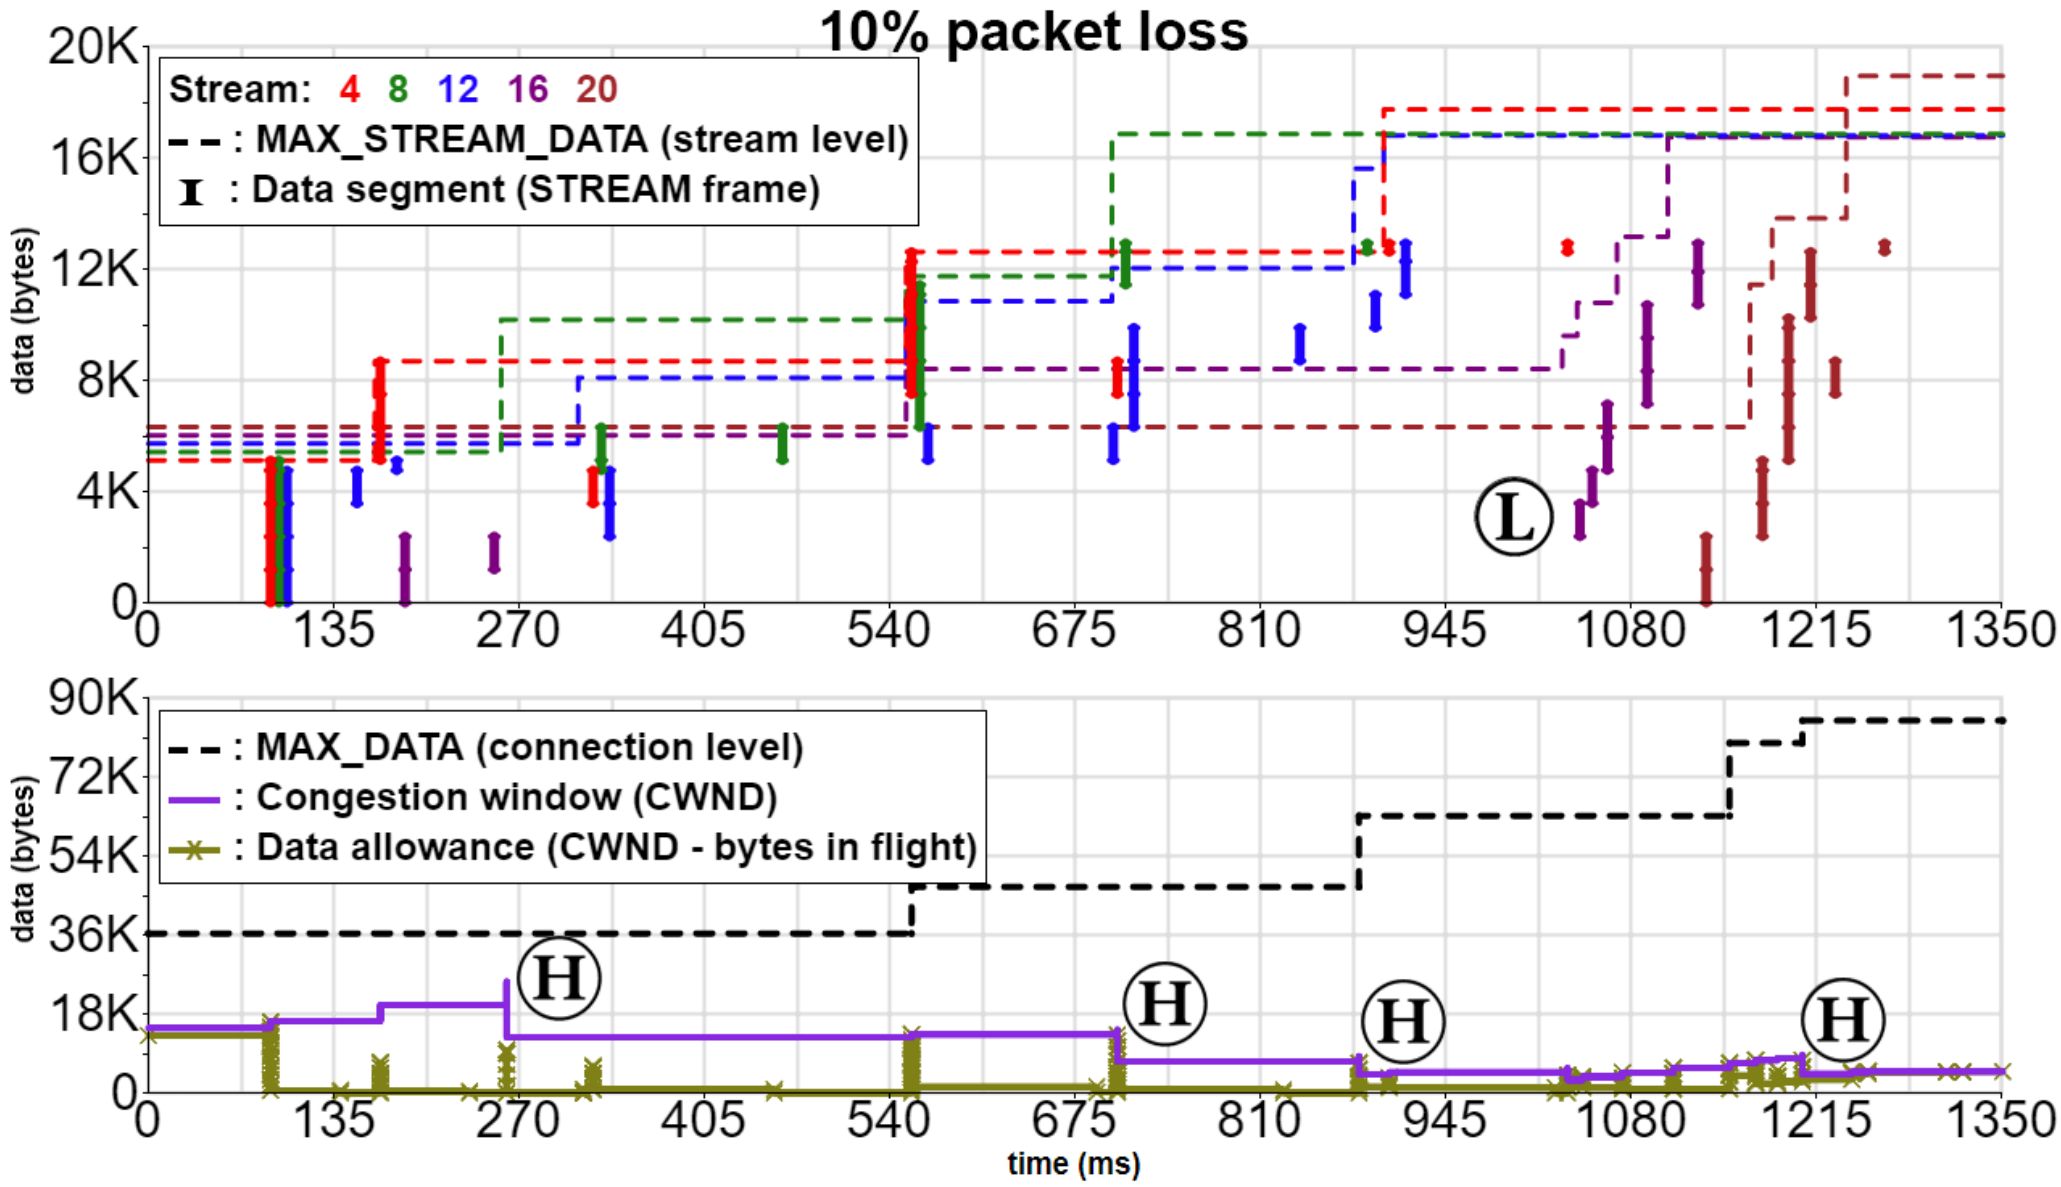

QUIC congestion control : quic-trace tool (Google)

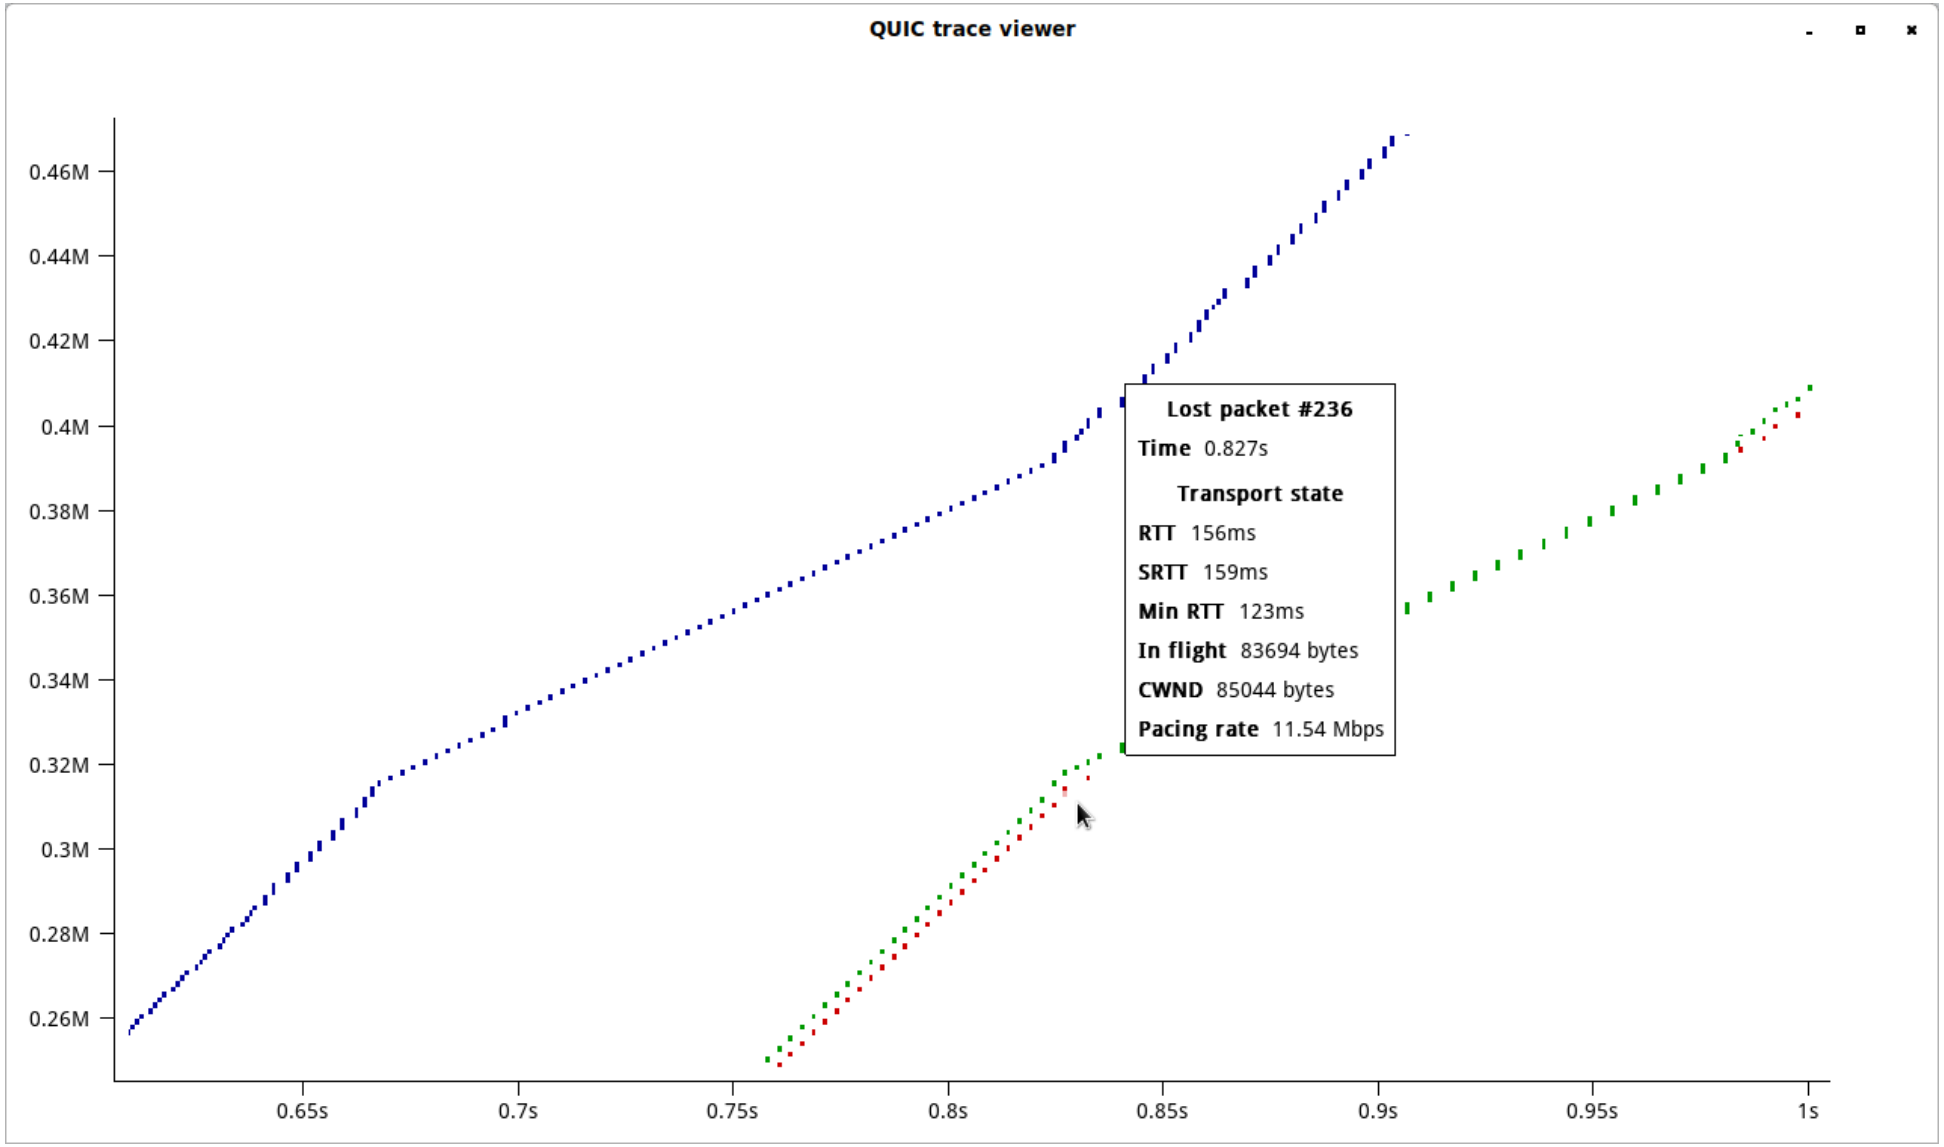

https://github.com/google/quic-trace

QUICvis

- POCs and demos (hardcoded data) available at
  - https://quic.edm.uhasselt.be
- Being reworked into a more flexible framework
  - Something usable in coming weeks
  - Prime-time ready by Prague (end of March)
  - V1: integration with quic-tracker

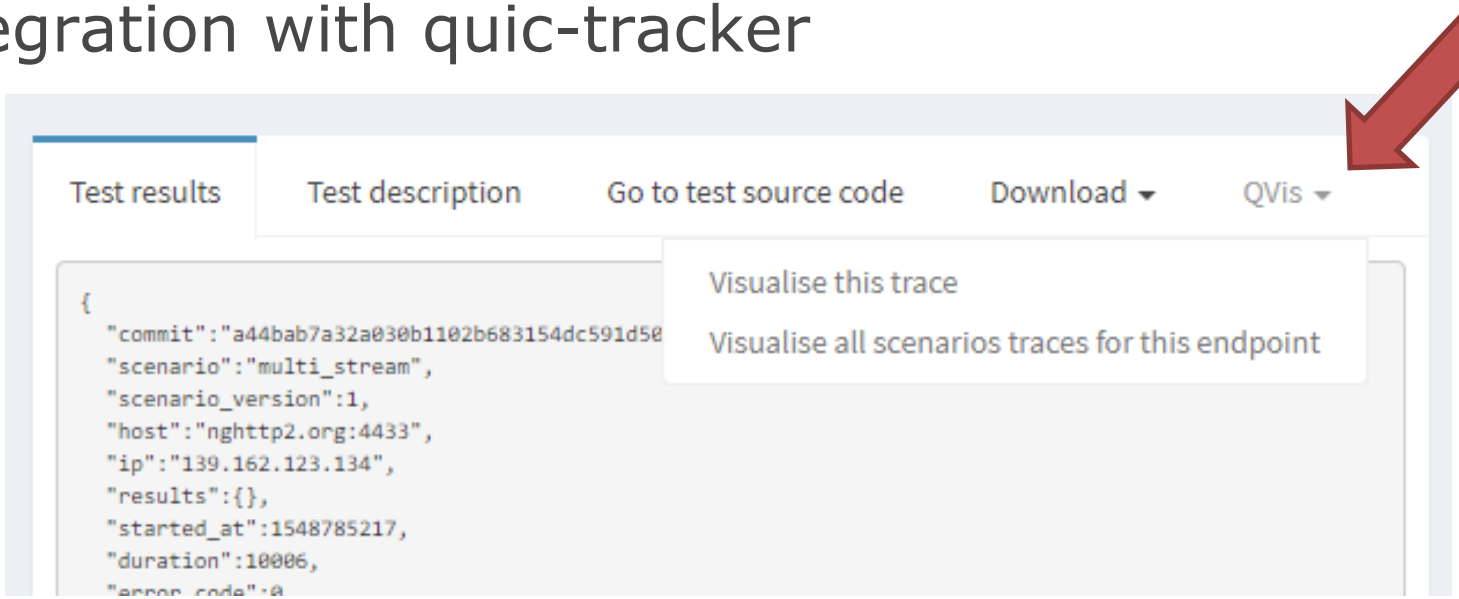

https://quic-tracker.info.ucl.ac.be/traces/20190129/229#

https://quicvis.edm.uhasselt.be:8443/#?list=https://quic-tracker.info.ucl.ac.be/traces/20190129/list/nghttp2.org:4433

QUICvis v1

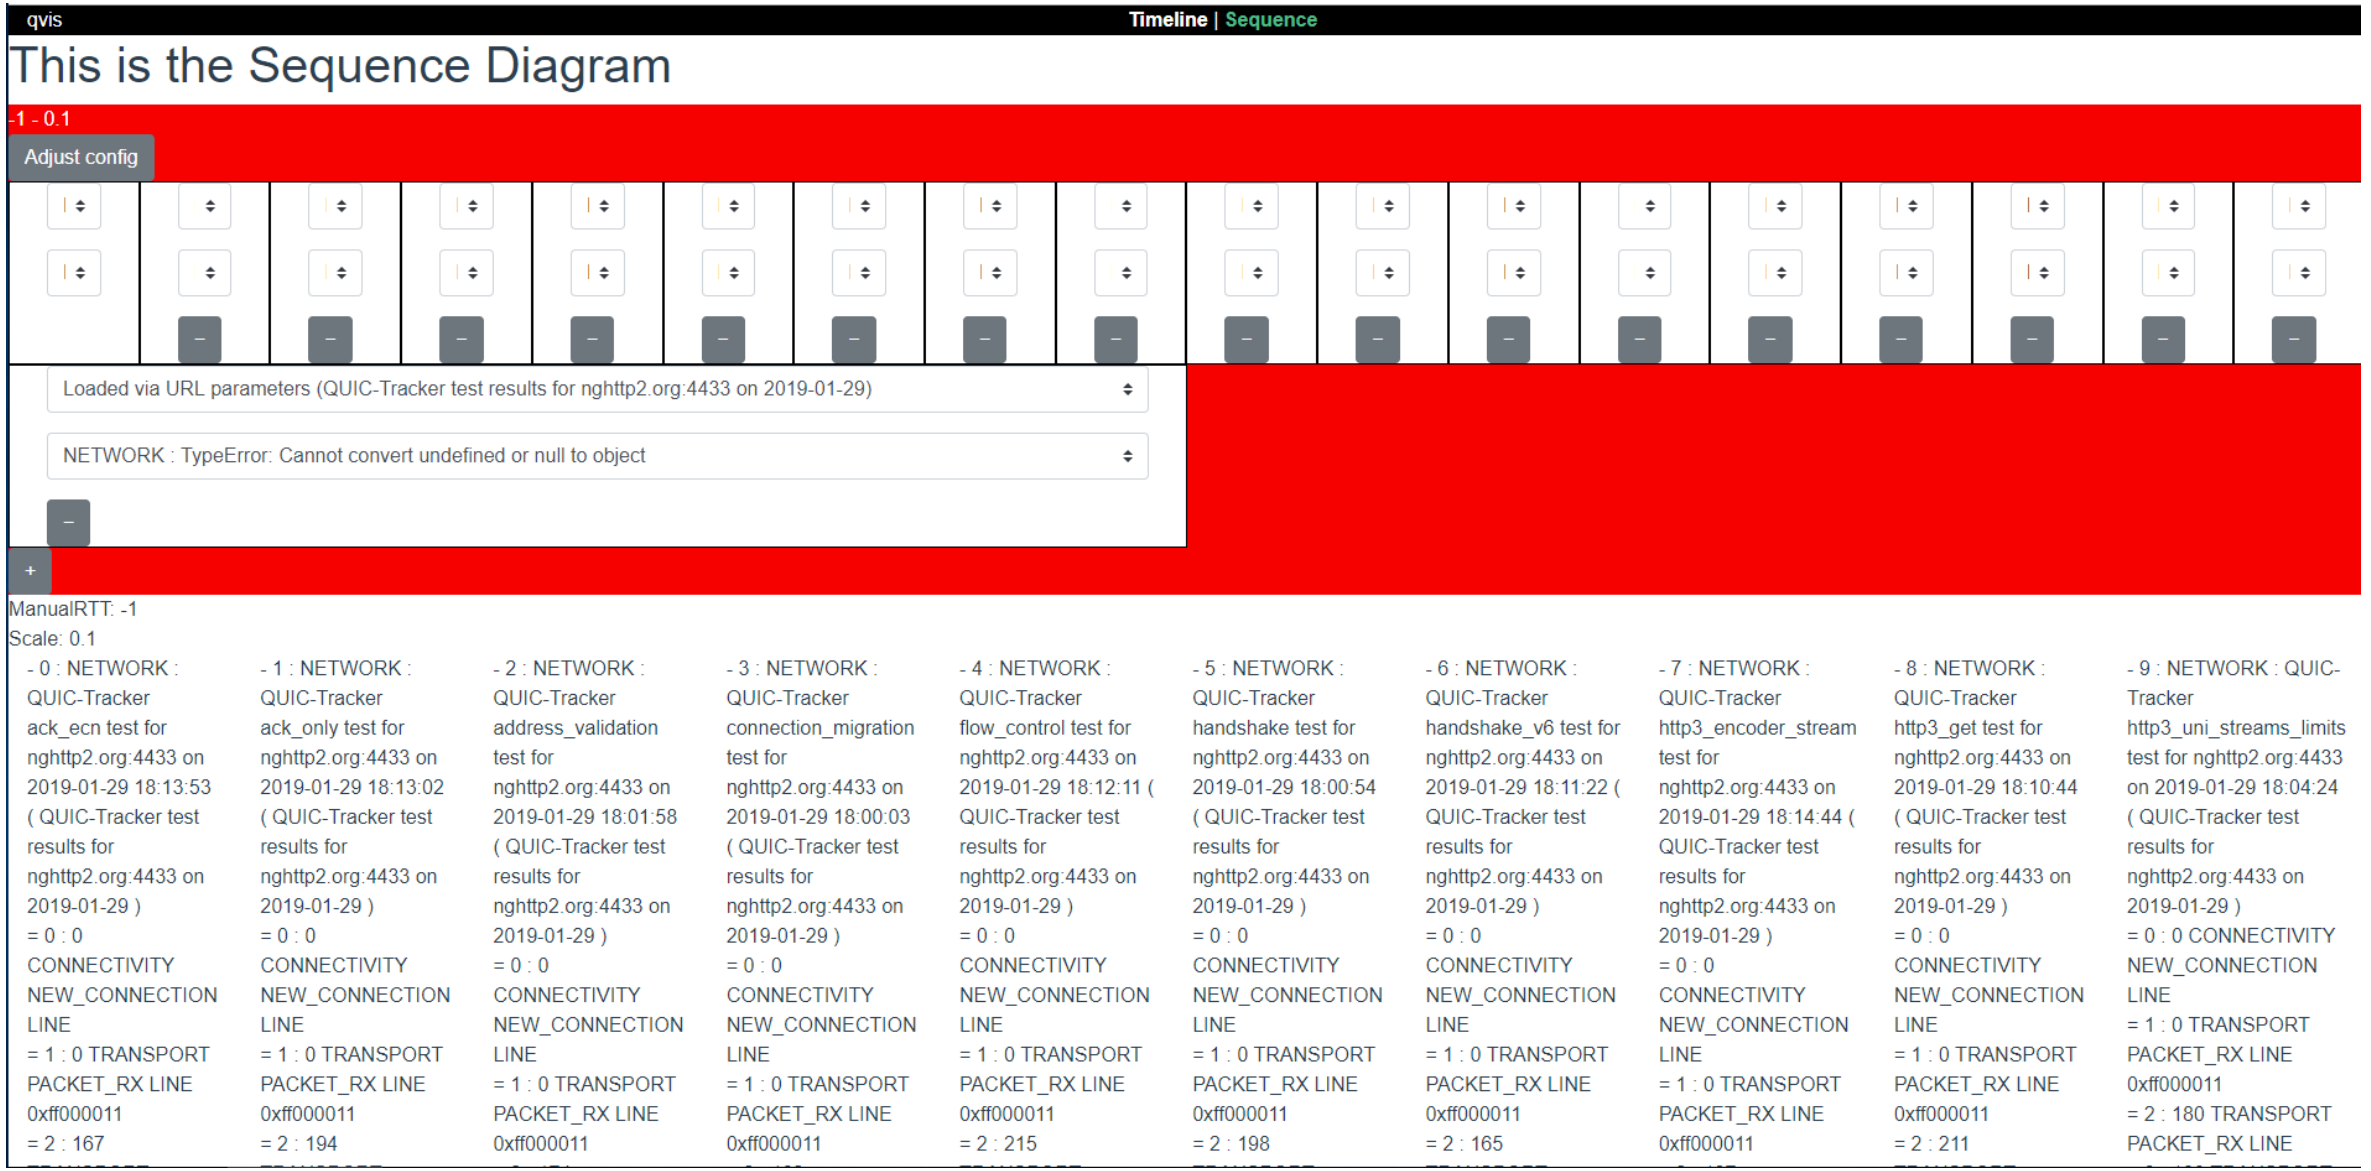
qvis
Timeline | Sequence
This is the Sequence Diagram
-1 - 0.1
Adjust config
Loaded via URL parameters (QUIC-Tracker test results for nghttp2.org:4433 on 2019-01-29)
NETWORK : TypeError: Cannot convert undefined or null to object
ManualRTT: -1
Scale: 0.1
- 0 : NETWORK : QUIC-Tracker ack_ecn test for nghttp2.org:4433 on 2019-01-29 18:13:53 (QUIC-Tracker test results for nghttp2.org:4433 on 2019-01-29) = 0 : 0 CONNECTIVITY NEW_CONNECTION LINE = 1 : 0 TRANSPORT PACKET_RX LINE 0xff000011 = 2 : 167
- 1 : NETWORK : QUIC-Tracker ack_only test for nghttp2.org:4433 on 2019-01-29 18:13:02 (QUIC-Tracker test results for nghttp2.org:4433 on 2019-01-29) = 0 : 0 CONNECTIVITY NEW_CONNECTION LINE = 1 : 0 TRANSPORT PACKET_RX LINE 0xff000011 = 2 : 194
- 2 : NETWORK : QUIC-Tracker address_validation test for nghttp2.org:4433 on 2019-01-29 18:01:58 (QUIC-Tracker test results for nghttp2.org:4433 on 2019-01-29) = 0 : 0 CONNECTIVITY NEW_CONNECTION LINE = 1 : 0 TRANSPORT PACKET_RX LINE 0xff000011
- 3 : NETWORK : QUIC-Tracker connection_migration test for nghttp2.org:4433 on 2019-01-29 18:00:03 (QUIC-Tracker test results for nghttp2.org:4433 on 2019-01-29) = 0 : 0 CONNECTIVITY NEW_CONNECTION LINE = 1 : 0 TRANSPORT PACKET_RX LINE 0xff000011
- 4 : NETWORK : QUIC-Tracker flow_control test for nghttp2.org:4433 on 2019-01-29 18:12:11 (QUIC-Tracker test results for nghttp2.org:4433 on 2019-01-29) = 0 : 0 CONNECTIVITY NEW_CONNECTION LINE = 1 : 0 TRANSPORT PACKET_RX LINE 0xff000011 = 2 : 215
- 5 : NETWORK : QUIC-Tracker handshake test for nghttp2.org:4433 on 2019-01-29 18:00:54 (QUIC-Tracker test results for nghttp2.org:4433 on 2019-01-29) = 0 : 0 CONNECTIVITY NEW_CONNECTION LINE = 1 : 0 TRANSPORT PACKET_RX LINE 0xff000011 = 2 : 198
- 6 : NETWORK : QUIC-Tracker handshake_v6 test for nghttp2.org:4433 on 2019-01-29 18:11:22 (QUIC-Tracker test results for nghttp2.org:4433 on 2019-01-29) = 0 : 0 CONNECTIVITY NEW_CONNECTION LINE = 1 : 0 TRANSPORT PACKET_RX LINE 0xff000011 = 2 : 165
- 7 : NETWORK : QUIC-Tracker http3_encoder_stream test for nghttp2.org:4433 on 2019-01-29 18:14:44 (QUIC-Tracker test results for nghttp2.org:4433 on 2019-01-29) = 0 : 0 CONNECTIVITY NEW_CONNECTION LINE = 1 : 0 TRANSPORT PACKET_RX LINE 0xff000011
- 8 : NETWORK : QUIC-Tracker http3_get test for nghttp2.org:4433 on 2019-01-29 18:10:44 (QUIC-Tracker test results for nghttp2.org:4433 on 2019-01-29) = 0 : 0 CONNECTIVITY NEW_CONNECTION LINE = 1 : 0 TRANSPORT PACKET_RX LINE 0xff000011 = 2 : 211
- 9 : NETWORK : QUIC-Tracker http3_uni_streams_limits test for nghttp2.org:4433 on 2019-01-29 18:04:24 (QUIC-Tracker test results for nghttp2.org:4433 on 2019-01-29) = 0 : 0 CONNECTIVITY NEW_CONNECTION LINE = 1 : 0 TRANSPORT PACKET_RX LINE 0xff000011 = 2 : 180 TRANSPORT PACKET_RX LINE

Logging format

- Current setup: transform .pcap to .json

- Problem: .pcap does not contain everything
  - Retransmit details (what was resent when and why?)
  - Exact latency information (Ack Delay on steroids)
  - HTTP/3 and QPACK state
  - Congestion control logic and state
  - *-> some can be inferred by heuristics, but not everything*

- “Solution”: Parse client/server logs from # implementations
  - Problem: extremely tedious and error-prone

Logging format

- The real solution: (de-facto) standard logging schema

- quic-trace has one
  - Focused on congestion-related events
  - Protocol-buffers

- We proposed one: qlog
  - Logs everything and more
  - .json-based (for easy integration with web-based tools)

qlog example : easy filtering and progressive enhancement

```
{"connectionid": "0x763f8eaf61aa3ffe84270c0644bdbd2b0d", "starttime": 1543917600,
 "fields":
    ["time","category",  "type",            "trigger",         "data"],
 "events": [
    [50,    "TLS",       "0RTT_KEY",        "PACKET_RX",       {"key": ...}],
    [51,    "HTTP",      "STREAM_OPEN",     "PUSH",            {"id": 0, "headers": ...}],
    ...
    [200,   "TRANSPORT", "PACKET_RX",       "STREAM",          {"nr": 50, "contents": "GET /ping.html", .
    [201,   "HTTP",      "STREAM_OPEN",     "GET",             {"id": 16, "headers": ...}],
    [201,   "TRANSPORT", "STREAMFRAME_NEW", "PACKET_RX",       {"id": 16, "contents": "pong", ...}],
    [201,   "TRANSPORT", "PACKET_NEW",      "PACKET_RX",       {"nr": 67, "frames": [16, ...], ...}],
    [203,   "RECOVERY",  "PACKET_QUEUED",   "CWND_EXCEEDED",   {"nr": 67, "cwnd": 14600, ...}],
    [250,   "TRANSPORT", "ACK_NEW",         "PACKET_RX",       {"nr": 51, "acked": 60, ...}],
    [251,   "RECOVERY",  "CWND_UPDATE",     "ACK_NEW",         {"nr": 51, "cwnd": 20780, ...}],
    [252,   "TRANSPORT", "PACKET_TX",       "CWND_UPDATE",     {"nr": 67, "frames": [16, ...], ...}],
    ...
    [1001,  "RECOVERY",  "LOSS_DETECTED",   "ACK_NEW",         {"nr": a, "frames": ...}],
    [2002,  "RECOVERY",  "PACKET_NEW",      "EARLY_RETRANS",   {"nr": x, "frames": ...}],
    [3003,  "RECOVERY",  "PACKET_NEW",      "TAIL_LOSS_PROBE", {"nr": y, "frames": ...}],
    [4004,  "RECOVERY",  "PACKET_NEW",      "TIMEOUT",         {"nr": z, "frames": ...}]
 ]}
```

qlog example : clear cause and effect

```
{"connectionid": "0x763f8eaf61aa3ffe84270c0644bdbd2b0d", "starttime": 1543917600,
 "fields":
    ["time","category",  "type",            "trigger",         "data"],
 "events": [
    [50,    "TLS",       "0RTT_KEY",        "PACKET_RX",       {"key":    }],
    [51,    "HTTP",      "STREAM_OPEN",     "PUSH",            {"id": 0, "headers": ...}],
    ...
    [200,   "TRANSPORT", "PACKET_RX",       "STREAM",          {"nr": 50, "contents": "GET /ping.html",
    [201,   "HTTP",      "STREAM_OPEN",     "GET",             {"id": 16, "headers": ...}],
    [201,   "TRANSPORT", "STREAMFRAME_NEW", "PACKET_RX",       {"id": 16, "contents": "pong", ...}],
    [201,   "TRANSPORT", "PACKET_NEW",      "PACKET_RX",       {"nr": 67, "frames": [16, ...], ...}],
    [203,   "RECOVERY",  "PACKET_QUEUED",   "CWND_EXCEEDED",   {"nr": 67, "cwnd": 14600, ...}],
    [250,   "TRANSPORT", "ACK_NEW",         "PACKET_RX",       {"nr": 51, "acked": 60, ...}],
    [251,   "RECOVERY",  "CWND_UPDATE",     "ACK_NEW",         {"nr": 51, "cwnd": 20780, ...}],
    [252,   "TRANSPORT", "PACKET_TX",       "CWND_UPDATE",     {"nr": 67, "frames": [16, ...], ...}],
    ...
    [1001,  "RECOVERY",  "LOSS_DETECTED",   "ACK_NEW",         {"nr": a, "frames":    }]
    [2002,  "RECOVERY",  "PACKET_NEW",      "EARLY_RETRANS",   {"nr": x, "frames": ...}],
    [3003,  "RECOVERY",  "PACKET_NEW",      "TAIL_LOSS_PROBE", {"nr": y, "frames": ...}],
    [4004,  "RECOVERY",  "PACKET_NEW",      "TIMEOUT",         {"nr": z, "frames": ...}]
 ]}
```

quic-trace example

```
enum EventType {
  UNKNOWN_EVENT = 0;
  PACKET_SENT = 1;
  PACKET_RECEIVED = 2;
  PACKET_LOST = 3;
  APPLICATION_LIMITED = 4;
  EXTERNAL_PARAMETERS = 5;
};

enum TransmissionReason {
  NORMAL_TRANSMISSION = 0;
  TAIL_LOSS_PROBE = 1;
  RTO_TRANSMISSION = 2;
  PROBING_TRANSMISSION = 3;
};
```

Triggers / Reasons

- Better than heuristics
- Can be tricky to implement (pass around additional state)

- Worth the effort?
- Not for everything?

Standard schema

- Options for integration
  - Directly implemented (maybe replaces current logging?)
  - Transformed from other format
    - E.g., Facebook database, Fastly also setting up logging infra
- Flexibility is important
  - Define minimal subset, rest is icing on the cake
  - You can support only the stuff you need (initially ☺)
  - Different levels of verbosity (debugging vs production)

Let us have some feedback

- By Prague
  - Tools and visualizations workable
  - Extensive logging schema (maybe as a draft?)
  - Results of people using the tools

  - Hopefully one of you also has qlog support built-in by then ☺

Let us have some feedback

- What will make you consider qlog?
  - Auto-validator?
  - Schema in certain format?
  - I send you a PR that integrates logging functions?
  - A full-blown logging library in C?
  - Also: what will make you dismiss it outright?
    - "json should burn in hell"
- Which features and use cases are most important to you?
  - Short term and longer term (E.g., debugging stack vs production tracing)
  - What would you use it for? What makes the extra work worth it?